

The Influence of Citation Formatting on Human Trust in Fabricated References

An Interdisciplinary Review at the Intersection of Information Science, Credibility Psychology, and Generative AI

Abstract

The proliferation of large language models (LLMs) as writing and research assistants has coincided with a well-documented tendency of these systems to generate citations that look authoritative but refer to no real work — a phenomenon termed "reference hallucination" or "citation fabrication." This paper synthesizes empirical and theoretical literature from information science, communication studies, cognitive psychology, and human–computer interaction to examine a specific and underexamined question: to what extent does the formatting of a citation — its style, completeness, apparent specificity, and surface conformity to disciplinary norms — shape human trust in a reference, independent of whether that reference actually exists? Drawing on cognitive heuristic models of credibility judgment (Metzger & Flanagin, 2013; Metzger, Flanagin, & Medders, 2010), the Modality-Agency-Interactivity-Navigability (MAIN) model of technology-cued credibility (Sundar, 2008), processing-fluency accounts of "truthiness" (Newman, Garry, Bernstein, Kantner, & Lindsay, 2012; Newman & Zhang, 2020), and recent empirical audits of citation fabrication in LLM output (Walters & Wilder, 2023; Linardon et al., 2025), we argue that citation format functions as a low-cost heuristic cue that is frequently mistaken for a validity signal. We review current detection and mitigation approaches, identify methodological and conceptual gaps in the literature — particularly the near-absence of controlled experiments that isolate formatting from content and source realism — and propose a research agenda for studying format-driven trust miscalibration in an era of AI-assisted writing.

Keywords: citation fabrication; reference hallucination; source credibility; cognitive heuristics; large language models; misinformation; truthiness; bibliographic verification; trust calibration; academic integrity

1. Introduction

Citations perform a peculiar dual function in scholarly and public discourse. On one hand, they are pointers — verifiable links to a specific, checkable artifact. On the other hand, they are rhetorical devices: their mere presence signals that a claim has been vetted, sourced, and is therefore more trustworthy than an unsupported assertion (Thornley et al., 2015). This second

function operates largely independently of the first. A reader encountering a properly formatted citation — complete with author names, a plausible title, a journal name, a year, and perhaps a Digital Object Identifier (DOI) — rarely verifies the reference before updating their belief in the claim it supports. This gap between the appearance of verifiability and the act of verification is precisely the space in which fabricated citations do their damage.

The urgency of this question has grown sharply with the mainstream adoption of generative AI tools for research and writing. Multiple empirical audits have found that LLMs such as GPT-3.5 and GPT-4 routinely generate references to papers, authors, and journals that do not exist, while nonetheless rendering them in impeccable citation style (Walters & Wilder, 2023). Walters and Wilder (2023) found that 55% of GPT-3.5-generated citations and 18% of GPT-4-generated citations in short literature reviews were entirely fabricated, and that even non-fabricated citations frequently contained factual errors — yet the vast majority nonetheless conformed closely to American Psychological Association (APA) formatting conventions. This last point is critical: formatting fidelity and factual fidelity are empirically dissociated in machine-generated references, and it is unknown how well human readers dissociate them perceptually.

This paper addresses that question by reviewing what is currently known — and, more importantly, what is not yet known — about the specific causal role of citation format (as distinct from citation content, source, or number) in shaping human trust judgments, with particular attention to the emerging problem of AI-generated fabricated references. We proceed as follows. Section 2 reviews the theoretical and empirical background on credibility heuristics, citation trust, and citation fabrication. Section 3 develops the central thematic argument: that citation format operates as a "fluency" and "authority" cue that is evaluated largely independently of underlying veracity. Section 4 surveys current technical and editorial approaches to detecting fabricated references. Section 5 discusses research challenges and limitations in the existing literature. Section 6 proposes a research agenda. Section 7 concludes.

2. Background and Related Work

2.1 Citations as Trust Signals in Human Communication

Long before generative AI, information scientists studied why researchers cite what they cite, and what citing behavior reveals about trust and authority. Thornley et al. (2015), in a qualitative study of researcher citation behavior, found that the decision to cite a source is frequently entangled with — but not identical to — a judgment of the source's trustworthiness and authority; researchers reported using citations both to signal engagement with a field and to lend borrowed credibility to their own claims. This dual signaling function is precisely what makes citations attractive targets for both legitimate persuasion and fabrication: a citation does rhetorical work even when the reader does not follow it.

Communication scholars have formalized this signaling process through models of heuristic information processing. Metzger and Flanagin (2013) synthesize a substantial body of research showing that, confronted with an overwhelming volume of online information, people rarely evaluate content systematically; instead they rely on a small set of cognitive heuristics — the reputation heuristic, the endorsement heuristic, the consistency heuristic, the

expectancy-violation heuristic, and, centrally for this paper, the authority heuristic — in which formal, official-looking, or citation-laden presentation is treated as a proxy for expertise and accuracy. Metzger, Flanagin, and Medders (2010), based on focus-group data from 109 participants, report that surface characteristics of a source — professional design, absence of errors, apparent thoroughness — were repeatedly and explicitly used by participants as fast, low-effort credibility cues, sometimes to the near-exclusion of any deeper verification. Bibliographic apparatus (footnotes, reference lists, DOIs) is a natural extension of this "professional design" cue: a densely and correctly formatted reference list looks like the product of careful, rule-following scholarship, regardless of whether every entry actually withstands scrutiny.

Sundar's (2008) Modality-Agency-Interactivity-Navigability (MAIN) model provides a complementary account rooted in human–computer interaction. The MAIN model argues that structural and technological affordances of an interface or artifact — not just its semantic content — cue heuristic judgments of credibility. A citation, in this framework, is itself a structural affordance: its typographic conventions (parenthetical author-year notation, italicized journal titles, hyperlinked DOIs) act as "authority cues" that trigger heuristic processing independent of whether the underlying agency (a real scholarly community) actually stands behind them. Extensions of the MAIN model to AI-mediated communication note that when the apparent "agent" behind a citation is ambiguous — as with an LLM assembling a reference list — users often default to treating machine output with either exaggerated skepticism or exaggerated authority, depending on other contextual cues, rather than applying consistent verification standards.

2.2 Empirical Work on Citations and Trust

A small but growing empirical literature has directly tested whether the mere presence of citations increases trust in AI-generated content, and how citation quality moderates that effect. Ding et al. (2025) conducted a randomized controlled experiment ($N = 303$ participants recruited via Prolific) in which participants submitted open-ended questions to a chatbot and received responses accompanied by zero, one, or five citations, which were either topically relevant or randomly substituted from other participants' sessions. The results are directly relevant to the present paper's thesis: the presence of any citation significantly increased self-reported trust in the response ($\beta = 0.394$, $p < .01$) relative to no citation at all — and this effect held even when the citations were random and irrelevant to the claim being supported, although random citations were trusted significantly less than topically valid ones ($\beta = -0.268$, $p < .01$). Critically, participants who actually moused over and inspected a citation subsequently reported lower trust than those who did not check, consistent with a "trust as anti-monitoring" account in which citation-checking is itself a marker of latent skepticism (Baier, 1986, as cited in Ding et al., 2025). Ding et al. also found no significant difference in trust between responses with one citation and those with five, suggesting a threshold or "satisficing" effect: a single citation is often sufficient to trigger the heuristic, after which additional citations yield diminishing returns.

While Ding et al. (2025) manipulated citation presence and relevance, they did not manipulate citation format — all citations in their study were rendered as simple numbered hyperlinks. This leaves open the question central to the present paper: whether increasing formatting elaborateness (full bibliographic detail, DOIs, journal names, page numbers) produces trust gains analogous to, independent of, or interactive with the relevance manipulation they studied.

2.3 The Scale of Machine-Generated Citation Fabrication

Parallel to this psychological literature, a rapidly expanding body of bibliometric and computational work has documented the prevalence of fabricated citations in LLM output. Walters and Wilder (2023), analyzing 636 citations across 84 GPT-3.5- and GPT-4-generated literature reviews on 42 multidisciplinary topics, found substantial fabrication rates that improved but did not disappear across model generations, and noted that fabricated citations frequently combined real author names with fictitious titles — a hybrid pattern that is particularly difficult for readers to detect through superficial inspection, since a familiar author name lends unearned credibility to an invented title. Linardon et al. (2025), publishing in *JMIR Mental Health*, extended this line of inquiry to a single discipline (mental health research), finding that GPT-4o's citation fabrication rate varied systematically with topic familiarity and prompt specificity: fabrication reached 19.9% of citations overall, with rates as high as 46% for specialized prompts on less publicly visible disorders (e.g., body dysmorphic disorder) and as low as roughly 6% for high-visibility topics (e.g., major depressive disorder). Among citations that did correspond to real publications, 45.4% contained bibliographic errors, most commonly incorrect or invalid DOIs — meaning that even "real" citations frequently fail formal verification.

The problem is not confined to informal or student use. A 2026 taxonomic analysis of accepted papers at the Conference on Neural Information Processing Systems (NeurIPS) — a top-tier, rigorously peer-reviewed AI venue — identified 100 hallucinated citations across 53 published papers (roughly 1% of accepted papers), despite review by three to five expert referees per submission, and classified the majority (66%) as "total fabrications" in which no corresponding work exists at all (Ansari, 2026). A related large-scale study of proceedings from four high-performance-computing conferences found a marked year-over-year increase in "mysterious citations" — references whose listed location contains no similar work — between 2021 and 2025 proceedings, coincident with the mainstream availability of generative AI writing tools (Anonymous, 2026). Journalistic coverage of a 2026 *Lancet* bibliometric study reported a roughly six-fold increase in fabricated citations appearing in published academic literature between 2023 and 2025 (STAT News, 2026). Together, these findings establish that fabricated but well-formatted citations are not a laboratory curiosity but an escalating feature of the actual scholarly and public information ecosystem — which sharply raises the practical stakes of understanding how format influences trust.

2.4 Processing Fluency and "Truthiness"

A final relevant strand of research comes from cognitive psychology's work on processing fluency and belief formation. Newman, Garry, Bernstein, Kantner, and Lindsay (2012) demonstrated that trivia claims (e.g., "the Mona Lisa has no eyebrows") are judged true significantly more often when accompanied by a related but strictly non-probative photograph — an image that provides no actual evidence for the claim — a phenomenon they termed "truthiness." Newman and Zhang (2020) review the mechanisms behind this effect, concluding that non-probative accompanying material increases the subjective ease, or fluency, with which a claim can be processed and mentally elaborated, and that this fluency is misattributed to the claim's truth value rather than to the decorative stimulus itself. Although this literature has focused on photographs, its underlying mechanism — non-probative, format-level cues inflating perceived truth via processing fluency — offers a directly transferable hypothesis for citation formatting: a full, properly styled reference (complete with volume, issue, page numbers, and DOI) may increase the fluency and "scholarly feel" of a claim in a manner functionally analogous to a decorative photograph, without providing

the reader any actual evidential content beyond what a bare, informal citation would provide.

3. Main Thematic Discussion: Format as a Heuristic Substitute for Verification

Synthesizing the literatures reviewed above, we advance a central argument: citation format operates as an authority and fluency heuristic that is processed largely independently of, and is frequently substituted for, verification of the citation's underlying existence and accuracy. This section unpacks the components of that argument.

3.1 Formatting as an Authority Cue

Under the MAIN model and the broader cognitive-heuristics tradition (Sundar, 2008; Metzger & Flanagin, 2013), authority is inferred not only from who purportedly produced content but from how that content is packaged. A citation rendered as "(Smith, 2019)" with a corresponding full entry — author, title, journal, volume, page range, DOI — presents multiple independent-seeming markers of institutional legitimacy. Each additional bibliographic field (a DOI, a page range, a named co-author) functions as what Metzger and Flanagin (2013) would classify under the "endorsement" and "consistency" heuristics: more detail is interpreted, consciously or not, as more corroboration, even though the fields are, in the case of a fabricated reference, mutually generated by the same ungrounded process and corroborate nothing. This is precisely the mechanism Walters and Wilder (2023) implicitly document when they note that fabricated ChatGPT citations nonetheless "adhere closely to APA format": the generative process learns to reproduce the syntax of scholarly authority without any binding to the semantics of a verifiable source.

3.2 The Dissociation Between Format Fidelity and Factual Fidelity

The empirical fabrication literature establishes a critical dissociation that any theory of format-driven trust must accommodate: formatting quality and factual accuracy vary largely independently in machine-generated references. Linardon et al. (2025) found that even genuine (non-fabricated) citations carried a 45.4% error rate, dominated by invalid DOIs — meaning a citation can look maximally verifiable (it has a DOI!) while that DOI resolves to nothing or to an unrelated work. Conversely, fabricated citations were frequently well-formed and stylistically indistinguishable from real ones at a glance. This dissociation means that a reader's felt sense of "this looks like a real, checkable reference" is a poor proxy for "this is a real, checkable reference" — and yet, per the heuristic-processing literature, felt sense is exactly what governs everyday trust judgments, because systematic, effortful verification (Metzger et al., 2010) is costly and is invoked only when stakes are high, suspicion is already aroused, or verification is unusually cheap.

3.3 Citation Presence, Random Content, and the Limits of the Heuristic

The Ding et al. (2025) experiment offers the clearest existing quantitative evidence that a bibliographic-apparatus effect operates even under adversarial conditions: trust rose significantly with the mere presence of a citation, even a provably irrelevant one, though not as much as with a relevant one, and dropped again once the participant actually clicked through to inspect it. This

pattern is consistent with a two-stage model: (1) an automatic, low-effort heuristic stage in which any citation-shaped object boosts trust, followed by (2) an optional, effortful verification stage that, when invoked, can partially or fully reverse the heuristic gain. Because format (not just presence) is a cue available entirely within stage 1 — a reader can register "this looks like a properly formatted academic citation" without clicking anything — we hypothesize, consistent with but not yet directly tested by existing data, that increasing formatting completeness would extend or amplify the stage-1 trust boost independent of relevance or existence, analogous to how a non-probative photograph amplifies truthiness independent of whether it depicts genuine evidence (Newman et al., 2012; Newman & Zhang, 2020).

3.4 Domain and Audience Moderators

The fabrication and trust literatures both suggest that this format effect is unlikely to be uniform across contexts. Linardon et al. (2025) show that fabrication rates — and by extension, the population of "format without substance" citations a reader is likely to encounter — vary by topic familiarity and prompt specificity, implying that specialist audiences reading about obscure subtopics may simultaneously face more fabricated citations and less domain knowledge with which to catch them, a troubling combination. Ding et al. (2025) additionally found that political and directly fact-checkable questions produced higher trust ratings independent of citation manipulation, suggesting that motivated reasoning and topic type interact with, rather than simply add to, citation-driven trust effects. Expertise is a plausible moderator as well: Metzger et al.'s (2010) focus-group participants reported falling back on heuristic, surface-level cues specifically because systematic evaluation was effortful, which implies that domain experts — for whom verification is comparatively cheap — should show a weaker format-trust coupling than lay readers, though this has not, to our knowledge, been directly tested for citations specifically.

4. Current Approaches and Developments

Responses to citation fabrication have thus far concentrated more on detection after the fact than on characterizing or correcting the human trust miscalibration that allows fabricated citations to cause harm before detection. Current approaches cluster into several categories.

Bibliographic verification tools and pipelines. A growing number of technical systems attempt to automatically verify LLM-generated citations against authoritative databases such as CrossRef, PubMed, and Semantic Scholar, flagging references whose metadata cannot be resolved or whose resolved metadata mismatches the claimed title, authors, or venue. Agentic frameworks for "traceable miscitation detection" (e.g., BibAgent-style systems referenced in recent LLM-hallucination audits) attempt to close the loop by cross-checking each generated reference at the point of composition rather than relying on post-hoc human review. These tools directly attack the dissociation problem described in Section 3.2 by re-coupling format to verifiable existence, but they presuppose that a verification step will actually be run — which, per the heuristic-processing literature, is precisely the step most human readers skip.

Editorial and publisher-level integrity measures. In response to documented increases in fabricated citations reaching even peer-reviewed venues (Ansari, 2026), publishers and conference organizers have begun adopting integrity-checking tools at the submission stage, and some journals now require authors to disclose and verify AI assistance in literature review.

The STAT News (2026) coverage of the Lancet bibliometric study explicitly frames this as an emerging institutional response to a measured six-fold rise in fabricated citations between 2023 and 2025, noting that publishers are turning to automated integrity tools as a countermeasure to what is described as "AI slop" entering the permanent scholarly record.

Model-level mitigation. On the generation side, retrieval-augmented generation (RAG) architectures attempt to ground citations in retrieved source documents rather than allowing the model to generate bibliographic fields from parametric memory, which Ding et al. (2025) note is intended to "improve the transparency and reliability" of AI output. However, Ding et al.'s own experimental design — presenting search-engine-retrieved but sometimes randomly mismatched citations — demonstrates that even grounded, non-hallucinated retrieval pipelines can still produce a citation-trust effect when a retrieved source is topically irrelevant, underscoring that "the citation is real" and "the citation actually supports this claim" are separate properties, both of which can fail independently of formatting quality.

Media-literacy and heuristic-awareness interventions. Building on the broader credibility-heuristics literature (Metzger & Flanagin, 2013), some information-literacy programs now explicitly teach the distinction between the appearance of sourcing and the substance of sourcing, encouraging spot-checking of citations (e.g., resolving DOIs, searching author names) as a low-cost countermeasure to the "trust as anti-monitoring" dynamic documented by Ding et al. (2025), in which the mere act of checking a citation reliably suppresses the unearned trust boost that an unchecked citation provides.

Notably absent from this landscape is any widely adopted intervention that targets the formatting-driven component of the trust effect specifically — for example, interface designs that visually flag unresolved DOIs at the point of reading, rather than only at the point of publication or generation.

5. Research Challenges and Limitations

Several conceptual and methodological limitations constrain the current evidence base and complicate straightforward inference from it.

Confounding of format, presence, and relevance. No study identified in this review manipulates citation formatting elaborateness (e.g., bare author-year vs. full APA entry with DOI vs. minimalist numbered bracket) as an independent factor while holding constant the citation's presence, relevance, and factual status. Ding et al.'s (2025) otherwise rigorous randomized design manipulates number and relevance/randomness of citations, but all citations were rendered in a single, uniform hyperlink format; consequently, the specific, isolated contribution of format to the trust effect they document remains inferred rather than directly measured. This is the central empirical gap this paper's argument depends on and the most important limitation of the present synthesis.

Reliance on self-reported trust. Both the citation-fabrication literature (e.g., Walters & Wilder, 2023; Linardon et al., 2025) and the citation-trust literature (Ding et al., 2025) rely substantially on either bibliometric ground-truth coding (for fabrication rates) or single-item self-reported trust scales (for trust effects); Ding et al. (2025) explicitly flag this as a limitation, noting that "different forms of trustworthiness may yield interesting results" and that behavioral measures of reliance

(e.g., whether a reader subsequently repeats or acts on a claim) may diverge from self-reported trust ratings, consistent with broader findings that increases in reported trust do not necessarily track increases in actual behavioral reliance on AI systems.

Sample and generalizability constraints. The Ding et al. (2025) sample, though randomized, was drawn from Prolific and skewed toward more technologically engaged, predominantly White (65%) participants, limiting confidence in generalizing citation-trust effects across demographic groups, education levels, or non-Western populations; the authors themselves note this and flag their "nonwhite" trust-difference finding as exploratory. Similarly, the fabrication-rate literature is dominated by English-language, APA- or numbered-style citation conventions; whether analogous formatting-driven trust effects generalize to other citation systems (e.g., Vancouver, Chicago, or non-Latin-script bibliographic traditions) is untested.

Model and topic specificity in fabrication studies. Linardon et al. (2025) explicitly note that their findings are specific to a single model (GPT-4o) and three mental-health topics, using comparatively simple prompts without advanced prompting or retrieval techniques; fabrication rates are known to vary substantially across model versions, and rapid model iteration means that quantitative fabrication-rate estimates risk becoming outdated even as the underlying formatting-trust dynamic likely persists.

Absence of longitudinal and downstream-consequence data. Existing work measures trust at a single point in time, immediately after exposure. Whether format-driven trust in a fabricated citation persists, decays, or compounds — for instance, through repeated exposure effects analogous to the illusory-truth effect, or through the citation being propagated into secondary sources — is unexamined in the reviewed literature, despite the STAT News (2026) report of citations "polluting the public record" over successive publication cycles, which implies a compounding, longitudinal dynamic that current single-shot experimental designs cannot capture.

Detection-tool evaluation gaps. Technical verification pipelines are generally evaluated on precision/recall against ground-truth fabrication labels, not on their downstream effect on human trust calibration; it is not established whether flagging an unresolved citation actually reduces reader trust proportionately, or whether users habituate to and begin ignoring such flags, a pattern well documented for other warning-label interventions in the misinformation literature.

6. Research Gaps and Future Directions

Building directly on the limitations identified above, we propose the following priority directions for future research.

1. **Factorial experiments isolating format from content and existence.** A minimal, high-value study would present participants with claims accompanied by (a) no citation, (b) a minimal citation (author-year only), (c) a moderately detailed citation (adds journal and year), and (d) a maximally detailed citation (adds volume, pages, and DOI), fully crossed with whether the citation is real, fabricated-but-plausible, or fabricated-with-a-broken-DOI. This design would directly estimate the marginal trust contribution of each additional bibliographic field, addressing the central gap identified in Section 5, and would allow a direct test of the "truthiness"-style fluency hypothesis advanced in Section 3.3 by

comparing citation-format effects to the established non-probative-photograph effect (Newman et al., 2012) within the same experimental paradigm.

2. Behavioral, not only self-reported, trust measures. Following Ding et al.'s (2025) own call for more robust trust measures, future work should incorporate behavioral indices — willingness to share a claim, willingness to cite the same reference in one's own writing, monetary or effort-based trust games, and eye-tracking or click-through measures of spontaneous verification behavior — to determine whether format-driven trust gains translate into consequential downstream reliance, or remain confined to self-report.

3. Cross-linguistic and cross-convention replication. Given that virtually all reviewed fabrication and trust studies operate within English-language, Western academic citation conventions (APA-style or numbered-bracket), replication across other citation systems and languages is needed to establish whether the formatting-trust effect is a general property of bibliographic apparatus or specific to conventions familiar to the study populations examined so far.

4. Expertise as a moderator. Direct tests of whether domain expertise attenuates format-driven trust (as predicted by heuristic-versus-systematic processing accounts; Metzger et al., 2010) versus whether experts are, perversely, more susceptible because they are more fluent readers of the format's surface conventions and less likely to click through, would help target literacy interventions appropriately — current evidence supports either hypothesis but tests neither directly for citations specifically.

5. Longitudinal and propagation studies. Research tracking whether a fabricated but well-formatted citation, once trusted and repeated, becomes self-reinforcing as it accumulates secondary citations of its own (an "citation laundering" dynamic implied but not directly measured by the NeurIPS and HPC-conference audits reviewed in Section 2.3) would connect individual-level trust psychology to the population-level "polluting the record" concern raised by STAT News (2026) and the Lancet bibliometric findings it reports.

6. Interface-level countermeasures evaluated for trust-calibration effects, not just detection accuracy. Future human–computer interaction research should evaluate whether real-time citation-verification indicators (e.g., a resolved/unresolved DOI badge shown at the point of reading) shift reader trust in the correct calibrated direction, and whether such indicators are subject to habituation or "warning fatigue" effects analogous to those documented for other credibility and misinformation labels.

7. LLM-user interaction design. Because RAG-based systems can still produce topically irrelevant but real citations (Ding et al., 2025), and because model-generation fabrication rates vary by topic familiarity (Linardon et al., 2025), future work should examine interface designs that make the relevance, not merely the existence, of a retrieved citation legible to users at a glance — potentially reducing reliance on format alone as a heuristic proxy for both existence and relevance simultaneously.

7. Conclusion

Across communication science, cognitive psychology, and the rapidly growing empirical literature on generative-AI citation fabrication, a consistent picture emerges: citations function as potent,

low-cost heuristic signals of credibility, and this signaling function is only loosely coupled to whether a citation is real, resolvable, or actually relevant to the claim it supports. Direct experimental evidence shows that trust rises with the mere presence of a citation, even a randomly mismatched one (Ding et al., 2025); parallel bibliometric evidence shows that machine-generated citations frequently achieve high formatting fidelity to disciplinary style guides while a substantial fraction are wholly or partially fabricated (Walters & Wilder, 2023; Linardon et al., 2025); and older but robust findings on non-probative cues and processing fluency (Newman et al., 2012; Newman & Zhang, 2020) supply a plausible cognitive mechanism — fluency misattributed to truth — for why a more elaborately formatted, but equally fabricated, citation might exert an even stronger unearned trust effect than a minimal one. What remains genuinely absent from the literature is a direct, factorial test of citation format as an isolated experimental variable. As AI-assisted writing tools become default infrastructure for both scholarly and lay information production, closing this gap is not merely of academic interest: it bears directly on how readers, reviewers, and institutions should design verification practices, interface warnings, and literacy interventions to prevent the surface conventions of scholarship from continuing to substitute for its substance.

Complete References

Ansari, S. (2026). Compound deception in elite peer review: A failure mode taxonomy of 100 fabricated citations at NeurIPS 2025 [Preprint]. arXiv.
<https://arxiv.org/abs/2602.05930>

Anonymous. (2026). The case of the mysterious citations [Preprint]. arXiv.

<https://arxiv.org/abs/2602.05867> Baier, A. (1986). Trust and antitrust. *Ethics*, 96(2), 231–260.

Ding, Y., Facciani, M., Poudel, A., Joyce, E., Aguinaga, S., Veeramani, B., Bhattacharya, S., & Weninger, T. (2025). Citations and trust in LLM generated answers [Preprint]. arXiv.
<https://arxiv.org/abs/2501.01303>

Linardon, J., Jarman, H. K., McClure, Z., Anderson, C., Liu, C., & Messer, M. (2025). Influence of topic familiarity and prompt specificity on citation fabrication in mental health research using large language models: Experimental study. *JMIR Mental Health*, 12(1), e80371.
<https://doi.org/10.2196/80371>

Metzger, M. J., & Flanagin, A. J. (2013). Credibility and trust of information in online environments: The use of cognitive heuristics. *Journal of Pragmatics*, 59, 210–220.
<https://doi.org/10.1016/j.pragma.2013.07.012>

Metzger, M. J., Flanagin, A. J., & Medders, R. B. (2010). Social and heuristic approaches to credibility evaluation online. *Journal of Communication*, 60(3), 413–439.
<https://doi.org/10.1111/j.1460-2466.2010.01488.x>

Newman, E. J., Garry, M., Bernstein, D. M., Kantner, J., & Lindsay, D. S. (2012). Nonprobative photographs (or words) inflate truthiness. *Psychonomic Bulletin & Review*, 19(5), 969–974.
<https://doi.org/10.3758/s13423-012-0292-0>

Newman, E. J., & Zhang, L. (2020). Truthiness: How non-probative photos shape belief. In R. Greifeneder, M. E. Jaffé, E. J. Newman, & N. Schwarz (Eds.), *The psychology of fake news: Accepting, sharing, and correcting misinformation* (pp. 90–114). Routledge/Taylor and Francis.

STAT News. (2026, May 7). Study finds explosion of "fraudulent" AI citations in academic papers. STAT.
<https://www.statnews.com/2026/05/07/lancet-study-finds-steep-rise-fraudulent-citations-academic-papers/>

Sundar, S. S. (2008). The MAIN model: A heuristic approach to understanding technology effects on credibility. In M. J. Metzger & A. J. Flanagin (Eds.), *Digital media, youth, and credibility* (pp. 73–100). MIT Press.

Thornley, C., Watkinson, A., Nicholas, D., Volentine, R., Jamali, H. R., Herman, E., Allard, S., Levine, K., & Tenopir, C. (2015). The role of trust and authority in the citation behaviour of researchers. *Information Research*, 20(3), Paper 677.
<http://informationr.net/ir/20-3/paper677.html>

Walters, W. H., & Wilder, E. I. (2023). Fabrication and errors in the bibliographic citations generated by ChatGPT. *Scientific Reports*, 13, Article 14045.
<https://doi.org/10.1038/s41598-023-41032-5>

Note on sourcing: This review draws exclusively on identifiable, publicly indexed scholarly and journalistic sources located via literature search. Two 2026 preprints are cited by title with arXiv identifiers where author attribution in the source was anonymized or unavailable at the time of writing; readers should consult the arXiv record directly for the most current author and venue metadata, as several of the AI-fabrication studies cited here describe a fast-moving and still-evolving empirical literature.